📗 Facilitator

Slide Deck Outline

Speaker Notes & Transitions — Day 1 & Day 2

Copilot Studio Workshop · Generated March 21, 2026

Slide Deck Outline

Day 1: Labs 00-12

Slide 1: Welcome and goals

- Welcome makers, IT pros, and developers into one shared Copilot Studio learning path.
- Set the expectation that Day 1 builds foundations and Day 2 turns them into an operational Hiring Agent solution.
- Define success as building, testing, evaluating, and explaining a production-ready direction. **Speaker note:** Open with outcomes and reassure the room that mixed roles are a strength, not a barrier.

Slide 2: Two-day outcomes

- Explain that Day 1 covers labs 00-12, Day 2 core shared path covers labs 13-24, and Lab 25 is the optional developer branch.
- Show how the workshop moves from fundamentals into governance, extensibility, evaluation, and optional developer workflow.
- Position every lab as a reusable pattern for real customer delivery, not a one-off demo. **Speaker note:** Use the full arc early so participants see why each lab matters beyond the classroom.

Slide 3: Mixed-audience roles

- Call out the maker lens: design the conversation, knowledge, prompts, cards, and flows.
- Call out the IT pro lens: secure the environment, channels, data access, and operational controls.
- Call out the developer lens: structure assets, extend behavior, and optionally use the VS Code workflow on Day 2. **Speaker note:** Frame the workshop as a shared delivery lifecycle where each role notices different risks and opportunities.

Slide 4: GA delivery guardrails

- Keep the core workshop on generally available capabilities and avoid non-GA dependencies.
- Use GPT-5 Chat as the recommended hands-on model and mention GPT-4.1 only when comparing platform context.
- Treat MCP as GA and use the in-product onboarding wizard as the only workshop setup path. **Speaker note:** Set these guardrails now so no later module drifts into unsupported or outdated guidance.

Slide 5: [TRANSITION] Into setup

- Move from workshop framing into environment readiness and UI confidence-building.
- Remind learners that early access issues are cheaper to solve before any build steps begin.
- Ask participants to surface blockers immediately instead of silently falling behind. **Speaker note:** This transition marks the shift from orientation to hands-on execution.

Slide 6: Lab Time: Lab 00

- Confirm everyone is signed in with the correct workshop account and environment.
- Use the timebox to verify browser readiness, maker access, and the facilitator support path.
- Tell participants that fast blocker reporting protects the rest of the day. **Speaker note:** Pause here until the room is truly ready to build from the same baseline.

Slide 7: Lab 00: Environment setup

- Open Copilot Studio, confirm the target environment, and review the landing experience.
- Verify that core menus, agent creation, and settings are visible for the workshop account.
- Capture any licensing or permission issues before moving on. **Speaker note:** Treat this lab as a go or no-go check for self-paced progress later in the day.

Slide 8: Module 1: Agents today

- Explain how Copilot Studio combines grounded answers, actions, channels, and operational controls.
- Distinguish lightweight question answering from true task completion and business process support.
- Connect platform value to the mixed audience in the room: adoption, governance, and extensibility. **Speaker note:** Keep the explanation practical and tied to business outcomes rather than product marketing language.

Slide 9: Lab Time: Lab 01

- Set the goal: explore the interface with a purpose instead of clicking every menu randomly.
- Ask learners to locate the surfaces they will reuse all workshop long.
- Encourage quick peer help so the room develops a common vocabulary. **Speaker note:** This lab reduces hesitation later by building shared UI familiarity now.

Slide 10: Lab 01: Intro to agents

- Open the main navigation, agent list, and starter creation path.
- Identify where overview, knowledge, topics, actions, publish, and test experiences live.

- Practice moving between those surfaces without losing orientation. **Speaker note:** Keep the lab lightweight so confidence rises before deeper authoring work begins.

Slide 11: Module 2: Studio foundations

- Map the key authoring surfaces to the agent lifecycle: define, ground, automate, test, publish, improve.
- Introduce the **Activity** page as the diagnostic habit participants should use throughout the workshop.
- Frame maintainability as a design choice that starts with structure, naming, and evidence. **Speaker note:** Give participants a mental model for how the workspace, diagnostics, and lifecycle fit together.

Slide 12: Activity page

- Use the **Activity** page to inspect what the agent saw, decided, and returned.
- Show how the view helps diagnose routing, grounding, and tool-selection issues without guesswork.
- Connect this practice to later evaluation, transcript analysis, and release decisions. **Speaker note:** Normalize evidence-based troubleshooting early so participants do not rely on intuition alone.

Slide 13: Diagnostics and maintainability

- Use the Microsoft architecture guidance to emphasize maintainable topics, clear ownership, and production diagnostics.
- Keep tool and knowledge names precise so orchestration decisions remain understandable over time.
- Design every workshop artifact as something a teammate could support after handoff. **Speaker note:** Tie well-architected thinking to everyday build choices instead of treating it as a separate review step.

Slide 14: Lab Time: Lab 02

- Set the objective: capture the core Copilot Studio fundamentals in the live UI.
- Ask learners to note which surface controls instructions, knowledge, actions, and publishing.
- Keep the room synchronized on one common baseline before branching into later patterns. **Speaker note:** Use this checkpoint to make sure foundational navigation is clear before more nuanced labs start.

Slide 15: Lab 02: Studio fundamentals

- Create or open a starter agent and review the core authoring surfaces in context.
- Confirm the overview, knowledge, action, publish, and diagnostic areas are easy to revisit.
- Record where the **Activity** page is accessed. **Speaker note:** The win condition is confident orientation, not exhaustive feature discovery.

Slide 16: Solution-aware delivery

- Explain why solutions matter for packaging assets, dependencies, and future lifecycle management.
- Separate workshop experimentation from the more durable structures teams need for reuse.
- Prepare participants to think in change sets and ownership boundaries, not isolated clicks. **Speaker note:** This framing helps mixed audiences see where delivery discipline starts to matter.

Slide 17: Lab Time: Lab 03

- Compare the speed of declarative agents with the control of more tailored designs.
- Ask learners to judge which scenarios fit a lighter approach and which do not.
- Keep the discussion anchored to business outcomes and governance needs. **Speaker note:** This lab works best when participants compare trade-offs, not winners.

Slide 18: Lab 03: Declarative agents

- Review a declarative agent pattern and note what can be delivered quickly.
- Compare the authoring effort with the limits on custom behavior, diagnostics, and extensibility.
- Capture one scenario where declarative design is enough and one where it is not. **Speaker note:** Use the exercise to sharpen judgment, not to create unnecessary complexity.

Slide 19: Lab Time: Lab 04

- Move the conversation from individual assets to solution packaging and lifecycle discipline.
- Tell learners to focus on structure, naming consistency, and what needs to travel together.
- Define success as understanding the delivery container, not just opening another screen. **Speaker note:** Keep the lab tightly scoped so it supports later operational thinking without slowing momentum.

Slide 20: Lab 04: Solutions

- Open or create a solution and review how related workshop artifacts are organized together.
- Confirm that the agent and its supporting assets can be reasoned about as a single deliverable.
- Note any environment dependencies that would matter in a real handoff or deployment. **Speaker note:** This lab introduces the packaging mindset participants will need once the

prototype becomes shared work.

Slide 21: Module 3: Reuse patterns

- Use prebuilt and reusable patterns to accelerate delivery without sacrificing clarity.
- Stress that reuse is valuable only when scenario fit, controls, and maintainability stay visible.
- Encourage teams to borrow proven structures rather than reinventing basics under time pressure. **Speaker note:** Position reuse as disciplined acceleration rather than shortcut-driven customization.

Slide 22: Lab Time: Lab 05

- Ask learners to inspect a prebuilt pattern quickly and judge what is portable to their scenario.
- Keep the focus on reusable ideas such as instructions, cards, topics, or flows.
- Discourage overcustomizing starter content during the workshop. **Speaker note:** The fastest way to learn here is to identify what should be reused and what should be rewritten.

Slide 23: Lab 05: Prebuilt agents

- Review a prebuilt agent or starter pattern and identify its strongest reusable pieces.
- Separate transferable structure from scenario-specific content that would need replacement.
- Capture one idea worth carrying forward and one gap that still requires custom design. **Speaker note:** Use the lab to build pattern recognition, not dependency on canned content.

Slide 24: Module 4: Custom agent design

- Shift from samples into owned design decisions around purpose, audience, and boundaries.
- Show that good custom agents feel intentional before they feel feature-rich.
- Keep maintainability visible by writing for future editors as well as current learners. **Speaker note:** This module marks the move from learning the product to designing with it.

Slide 25: Architecture for grounded agents

- Align instructions with the actual tools and knowledge the agent can use so behavior stays trustworthy.
- Prefer smaller maintainable building blocks over giant monolithic topics or prompts.
- Plan diagnostics, support ownership, and production readiness alongside the first build. **Speaker note:** Use the architecture references to reinforce that grounded, supportable agents are designed deliberately.

Slide 26: Lab Time: Lab 06

- Set the deliverable: a clearly named custom agent with a crisp purpose statement.
- Ask learners to keep the first version simple, testable, and easy to explain.
- Use consistent naming so facilitator support stays efficient. **Speaker note:** The right outcome here is a strong foundation, not maximum scope.

Slide 27: Lab 06: Custom agent

- Create or refine the workshop custom agent in the target environment.
- Add initial instructions and connect the first knowledge source or placeholder.
- Test whether the agent purpose is immediately understandable in the UI and chat. **Speaker note:** Coach participants to make the first version understandable before they make it sophisticated.

Slide 28: Module 5: Topic design

- Use topic guidance to keep entry points clear, reusable, and low overlap.
- Explain how bite-size topics and redirects improve maintainability over time.
- Warn that ambiguous triggers create support cost long before they create visible bugs. **Speaker note:** Bring the topic authoring best practices into the room as a maintainability story, not just a design preference.

Slide 29: Bite-size reusable topics

- Split large conversation logic into smaller reusable topics with clear intent.
- Watch multiple-topic matches as a signal that trigger phrases overlap too much.
- Use redirects and shared endings instead of copying the same logic repeatedly. **Speaker note:** Show that topic structure is one of the simplest ways to improve both quality and maintenance.

Slide 30: Lab Time: Lab 07

- Set the task: build one focused topic path before adding optional branches.
- Ask learners to test realistic trigger wording instead of idealized commands.
- Make the done state a clean start, middle, and exit that others can follow. **Speaker note:** A small reliable topic teaches more than a large unfinished one.

Slide 31: Lab 07: Topics and triggers

- Create a topic with clear trigger phrases, required variables, and a clean exit.
- Test the trigger against realistic phrasing and watch for overlap or confusion.
- Record one improvement that would make the topic easier to maintain later. **Speaker note:** Keep the exercise grounded in clarity, because clarity is what scales.

Slide 32: Adaptive cards that travel

- Use Adaptive Cards to validate input, gather structured data, and enrich the conversation experience.
- Design cards with host compatibility, accessibility, and simple support expectations in mind.
- Keep cards purposeful so they improve the workflow instead of becoming decoration. **Speaker note:** Anchor card design to supported host behavior so the workshop stays realistic and reusable.

Slide 33: Lab Time: Lab 08

- Ask learners to add one card only where it genuinely improves data quality or clarity.
- Collect only the inputs needed by the next business step.
- Keep the implementation simple enough to test within the timebox. **Speaker note:** Use the lab to show how UI can strengthen a conversation without overwhelming it.

Slide 34: Lab 08: Adaptive cards

- Add or refine an Adaptive Card interaction inside the workshop agent.
- Test the card in the supported host context and verify the values are captured cleanly.
- Confirm the next prompt, topic, or action benefits from the structured input. **Speaker note:** Treat the card as a precision tool for the workflow, not as a generic form.

Slide 35: Module 6: Actions and events

- Connect conversation to business process only when the scenario needs a side effect or follow-up action.
- Explain why actions and events should be explicit, auditable, and easy to diagnose.
- Prepare the room to think about orchestration as workflow design rather than freeform prompting. **Speaker note:** This module is where the agent starts to become part of a business system.

Slide 36: Generative orchestration rules

- Write instructions that constrain outputs, tool choices, and response format for business use.
- Keep prompts grounded in configured tools and knowledge rather than wishful instructions.
- Prefer repeatable, useful outputs over creative variance when the workflow needs consistency. **Speaker note:** Use the generative-mode guidance to show that better instructions reduce ambiguity everywhere else.

Slide 37: Lab Time: Lab 09

- Set the goal: wire one meaningful action with visible inputs and outputs.
- Tell learners to start with the smallest useful scope before adding branches or exceptions.
- Define success as a testable business step, not just a connected flow icon. **Speaker note:** Slow down slightly here because action wiring is a common confidence dip.

Slide 38: Lab 09: Agent flows

- Connect the agent to a flow or action that uses captured context.
- Inspect what is passed in, what returns, and what the user sees on success or failure.
- Note any permission or dependency concerns that would matter after the workshop. **Speaker note:** Use the lab to show why automation design needs both technical and business clarity.

Slide 39: Lab Time: Lab 10

- Move from user-initiated chat into event-driven automation patterns.
- Ask learners to decide what evidence proves the trigger fired correctly.
- Keep the scenario small enough that the room can reason about the outcome quickly. **Speaker note:** This lab helps participants think beyond the test pane and into operational workflows.

Slide 40: Lab 10: Event triggers

- Configure or review an event-triggered path in the workshop scenario.
- Validate the trigger condition, downstream action, and user-facing evidence.
- Capture what would need monitoring or admin review before broader rollout. **Speaker note:** Tie the trigger path back to observable operations rather than invisible automation magic.

Slide 41: [TRANSITION] From build to publish

- Shift from building the prototype to reviewing release readiness and ownership.
- Bring together naming, access, diagnostics, channels, and governance in one conversation.
- Set the expectation that publish readiness is broader than a single button. **Speaker note:** Use this transition to refocus the room on operational credibility.

Slide 42: Lab Time: Lab 11

- Set the task: inspect publish readiness, supported channels, and likely admin dependencies.
- Ask learners to note questions about audience, branding, and support ownership.
- Keep the room aligned on readiness signals instead of forcing identical publish outcomes. **Speaker note:** This checkpoint works best when participants think like future owners, not just workshop attendees.

Slide 43: Lab 11: Publish agent

- Review the publish flow and channel options that matter for enterprise rollout.
- Compare readiness considerations for Teams, Microsoft 365 Copilot, website, and later channel expansion.
- Capture any policy, branding, or support blockers that would delay go-live. **Speaker note:** Reinforce that publishing is a coordination task across makers, admins, and business owners.

Slide 44: Lab Time: Lab 12

- Close Day 1 by switching from technical readiness into commercial and adoption readiness.
- Ask learners to connect workshop actions to credits, licensing, and value measurement.
- Invite IT pros and sponsors to listen for post-launch planning signals. **Speaker note:** This lab gives the room the vocabulary to talk about scale, cost, and value responsibly.

Slide 45: Lab 12: Licensing and ROI

- Map workshop scenarios to Copilot Studio usage, licensing decisions, and operational planning.
- Explain when advanced actions, external channels, or automation consume Copilot Studio capacity.
- Introduce ROI analytics as part of post-launch measurement, not an afterthought. **Speaker note:** End Day 1 by linking technical progress to commercial credibility and Day 2 readiness.

Day 2: Core Labs 13-24 + Optional Lab 25

Slide 46: [TRANSITION] Day 2 restart

- Reconnect the room to the Day 1 baseline and carry the strongest practices forward.
- Explain that Day 2 adds orchestration, safety, enterprise data, evaluation, and an optional developer workflow after the core path.
- Set the expectation that every module now emphasizes production readiness more explicitly. **Speaker note:** Use the restart to create momentum without replaying the entire previous day.

Slide 47: Module 7: Hiring architecture

- Anchor the day around a Hiring Agent solution that feels closer to real operational complexity.
- Show the maintainable architecture: core agent, delegated capabilities, governed data, actions, and diagnostics.
- Align makers, IT pros, and developers around shared delivery responsibilities. **Speaker note:** Position the Hiring Agent as a structured solution, not a larger prompt.

Slide 48: Multi-agent operating model

- Use clear boundaries so each agent or capability owns a focused responsibility.
- Explain how handoffs, escalation, and ownership reduce confusion when workflows expand.
- Prefer coordinated components over one oversized agent that is hard to test or support. **Speaker note:** This framing helps the room understand why structured orchestration improves maintainability.

Slide 49: Lab Time: Lab 13

- Set up the Hiring Agent baseline with the correct environment, naming, and scenario context.
- Ask learners to confirm data and knowledge prerequisites before changing behavior.
- Define success as a reliable starting point for the rest of Day 2. **Speaker note:** A stable baseline now prevents compounded confusion in later connected labs.

Slide 50: Lab 13: Hiring setup

- Create or open the Hiring Agent and confirm the main authoring surfaces are ready.
- Review the scenario goals, dependencies, and any pre-staged assets the class will reuse.
- Capture what must stay consistent across later labs for testing to remain meaningful. **Speaker note:** Use the setup lab to align expectations before the day becomes more complex.

Slide 51: Lab Time: Lab 14

- Strengthen the instructions before adding more tools, data, or automation.
- Ask learners to write for maintainability, scope clarity, and recruiter-safe outputs.
- Keep the first improvement easy to test in the chat experience. **Speaker note:** Instructions are the control plane for later behavior, so invest here before extending further.

Slide 52: Lab 14: Agent instructions

- Refine the Hiring Agent instructions with clear scope, tone, escalation, and formatting guidance.
- Align the instructions to configured tools and knowledge instead of asking the model to invent capability.
- Test one direct request and one ambiguous request to see whether the guidance holds. **Speaker note:** This lab is most valuable when participants connect instruction quality to predictable outcomes.

Slide 53: Lab Time: Lab 15

- Move into multi-agent collaboration with a clear purpose for each delegated capability.
- Ask teams to keep each role narrow enough that ownership stays understandable.
- Define success as a handoff model that is easy to explain and support. **Speaker note:** Keep the collaboration pattern simple so the operating model stays visible.

Slide 54: Lab 15: Multi-agent team

- Configure or review a multi-agent hiring pattern with clear delegation boundaries.
- Confirm each specialized capability has a purpose, entry point, and expected outcome.
- Validate that the coordination logic stays understandable in the UI and test flow. **Speaker note:** The main lesson is disciplined orchestration, not maximum branching.

Slide 55: Module 8: Automation and models

- Add automation and model choices without losing operational control.
- Compare models by task fit, response shape, cost, and consistency rather than trend.
- Keep the workshop baseline on GA capabilities so the room can reproduce outcomes later. **Speaker note:** Frame model choice as an implementation decision that should stay tied to task quality.

Slide 56: Lab Time: Lab 16

- Automate one useful hiring workflow step with a clear trigger and outcome.
- Tell learners to keep the scenario observable and explainable to nontechnical stakeholders.
- Capture what success evidence should appear if the automation works. **Speaker note:** Use the timebox to keep automation grounded in one measurable business step.

Slide 57: Lab 16: Trigger automation

- Configure trigger-based behavior tied to a concrete hiring event or process step.
- Confirm the event path, connection state, and resulting action are visible in the UI.
- Record any dependency that would matter for tenant readiness after the workshop. **Speaker note:** This lab teaches participants to treat automation as an operational commitment, not just a feature.

Slide 58: Lab Time: Lab 17

- Compare models using the same hiring task so the differences stay meaningful.
- Use GPT-5 Chat as the recommended hands-on baseline for the workshop.
- Mention GPT-4.1 only as platform context when participants need a comparison point. **Speaker note:** Keep the comparison lightweight so the room focuses on fit, not benchmarking theater.

Slide 59: Lab 17: Model selection

- Run the comparison in the UI and judge structure, speed, and usefulness for hiring tasks.
- Select the model that best balances reliable output with the workshop scenario needs.

- Save the chosen baseline so later labs evaluate against the same model. **Speaker note:** Model selection matters most when the room can explain why a choice was made.

Slide 60: Safety by layered controls

- Combine AI disclosure, instructions, moderation, error handling, and testing into one safety model.
- Use per-prompt content moderation (Low/Moderate/High slider) where a specific prompt or tool needs tighter moderation than the rest of the agent.
- Treat safety as something to validate with evidence, not something to assume after configuration. **Speaker note:** This slide sets up content moderation as a layered operational control rather than a toggle.

Slide 61: Lab Time: Lab 18

- Harden the Hiring Agent with realistic safe-use guardrails.
- Ask learners to note which protection layer responds to each risky prompt.
- Keep one clean in-scope prompt as a control so safe behavior remains useful behavior. **Speaker note:** Red-team style testing works best when the room compares outcomes across the same prompt set.

Slide 62: Lab 18: Content moderation

- Set the moderation posture and review how refusal or redirect messaging appears to users.
- Inspect prompt-level or tool-level sensitivity only where targeted tightening is justified.
- Verify the agent refuses unsafe, irrelevant, or policy-breaking requests while staying helpful in scope. **Speaker note:** Emphasize that precise controls beat blanket loosening or blanket blocking.

Slide 63: Module 9: Multimodal and data

- Add richer inputs only when they improve the hiring workflow and decision quality.
- Connect the conversation to governed business data without losing traceability.
- Keep every output reviewable, role-appropriate, and easy to justify. **Speaker note:** This module keeps capability growth tied to business value and control.

Slide 64: Multimodal prompt framing

- Decide what a document or image should contribute to the hiring workflow before prompting on it.
- Constrain extraction and summarization to job-related outcomes the room can validate.

- Use clarifying follow-up when the input is incomplete or ambiguous. **Speaker note:** Well-framed multimodal prompts stay useful because they are scoped to a real decision.

Slide 65: Lab Time: Lab 19

- Test one multimodal path with a controlled document or image input.
- Ask learners to verify both usefulness and the need for human review.
- Keep the exercise grounded in the Hiring Agent scenario rather than novelty. **Speaker note:** Use the lab to show where multimodal input helps and where human judgment still matters.

Slide 66: Lab 19: Multimodal prompts

- Run a multimodal prompt workflow in the UI and inspect the returned summary or extraction.
- Confirm the output remains recruiter-friendly, job-related, and easy to review.
- Note one improvement that would make the workflow safer or more maintainable. **Speaker note:** The value here comes from controlled usefulness, not from using every possible input type.

Slide 67: Dataverse grounding patterns

- Use Dataverse when the scenario needs governed operational data, structured records, and role-based access.
- Align table visibility, permissions, and traceability before trusting the experience in front of users.
- Connect data-grounded behavior to later diagnostics, transcript analysis, and improvement cycles. **Speaker note:** Tie Dataverse readiness to both business value and operational accountability.

Slide 68: Lab Time: Lab 20

- Connect the Hiring Agent to Dataverse safely and verify only the needed data is in scope.
- Ask learners to validate permissions early so the room does not mistake access issues for design issues.
- Define success as visible, relevant, governed data access in the UI. **Speaker note:** Treat this as a readiness checkpoint because data access variability is common across tenants.

Slide 69: Lab 20: Dataverse grounding

- Open or configure the Dataverse grounding path and confirm the target data is visible.
- Test a hiring question that benefits from governed business data instead of general knowledge only.

- Capture any mismatch between scenario needs, data quality, and available permissions. **Speaker note:** Keep the lab focused on governed usefulness rather than deep schema exploration.

Slide 70: Document outputs with guardrails

- Generated documents need clear inputs, approved templates, and obvious review points.
- Focus on repeatable hiring artifacts such as interview packs, summaries, or recruiter briefs.
- Make output quality measurable before anyone treats the result as production-ready. **Speaker note:** This framing helps document generation feel operational rather than magical.

Slide 71: Lab Time: Lab 21

- Generate one document-oriented output only after the required context is present.
- Ask learners to review the result for structure, policy fit, and next-step usefulness.
- Use the timebox to document one improvement instead of endlessly polishing wording. **Speaker note:** A good first output teaches more than a perfect tenth revision.

Slide 72: Lab 21: Document generation

- Create or review the document-generation step tied to the Hiring Agent scenario.
- Validate the produced artifact against the expected structure, tone, and job-related scope.
- Record one guardrail or template improvement before considering wider use. **Speaker note:** Keep the result tied to a business artifact someone could actually review or use.

Slide 73: [TRANSITION] Into MCP

- Move from internal orchestration into governed extensibility with external tools.
- Keep access narrow, auditable, and easy to explain to admins and sponsors.
- Use only the supported GA setup path so the workshop remains reproducible. **Speaker note:** This transition sets the expectation that extensibility must stay governed to stay credible.

Slide 74: Module 10: MCP and extensibility

- Treat MCP as a generally available capability in this workshop.
- Use the in-product onboarding wizard to add or review MCP servers inside Copilot Studio.
- Keep tool scope narrow, auditable, and explainable to admins; avoid non-GA detours.
- Avoid fallback branching, manual secret editing, or any non-wizard setup path. **Speaker note:** State the supported path clearly so the room does not anchor on outdated guidance and extensibility stays governed.

Slide 75: Lab Time: Lab 22

- Add MCP through the supported UI path and keep the scope limited to one useful hiring task.
- Ask learners to validate why the server belongs in the scenario before they enable it.
- Define success as a governed tool connection that is visible and explainable. **Speaker note:** Use this lab to connect extensibility to security and supportability, not just capability.

Slide 76: Lab 22: MCP integration

- Open the MCP onboarding wizard and add or review a supported server in the agent.
- Confirm the tool appears with the right connection state and purpose in the UI.
- Test one safe, auditable MCP-assisted task that fits the Hiring Agent scenario. **Speaker note:** Keep the example small so the room learns the pattern without overextending scope.

Slide 77: Module 11: Channels and feedback

- Close the operational loop by pairing publishing readiness with user feedback and support expectations.
- Match channel choices to audience, policy, and ownership rather than novelty.
- Turn feedback into evidence that can influence backlog, evaluation, and release decisions. **Speaker note:** This module connects user reach with the quality loop that keeps the agent improving.

Slide 78: Channel overview: Teams and more

- Compare Teams, Microsoft 365 Copilot, WhatsApp, and web experiences from a readiness and governance perspective.
- Call out how branding, support, compliance, and channel policy affect rollout decisions.
- Explain why channel choice changes monitoring, adoption, and change-control expectations. **Speaker note:** Use the channel overview to broaden the conversation from publishing mechanics to operating reality.

Slide 79: Lab Time: Lab 23

- Capture user feedback intentionally instead of treating comments as informal noise.
- Ask learners to decide who reviews feedback, how often, and what action follows.
- Keep the feedback path tied to quality improvement and adoption learning. **Speaker note:** This lab turns user reaction into something the team can operationalize.

Slide 80: Lab 23: User feedback

- Enable or review the feedback capture path and verify where feedback lands.
- Test how a user can submit feedback after a conversation or outcome.

- Connect feedback signals to future tuning, release decisions, and support processes. **Speaker note:** Show that feedback is most useful when it has an owner and a follow-up path.

Slide 81: Module 12: Evaluation and ROI

- Use evaluation as a repeatable release gate instead of relying on ad hoc chat testing.
- Pair evaluation evidence with transcript analytics and ROI analytics to tell a fuller value story.
- Connect quality results to rollout, staffing, and investment decisions after the workshop. **Speaker note:** This final module brings quality, insight, and business value into one decision framework.

Slide 82: Lab Time: Lab 24

- Run a repeatable quality review with scored evidence instead of intuition alone.
- Ask learners to compare pass rate, failure reasons, and next-step fixes.
- Use the activity map and transcript evidence to prepare one improvement loop. **Speaker note:** This lab matters most when participants leave with a reusable QA habit, not just one run.

Slide 83: Lab 24: Evaluation and QA

- Create or run an evaluation test set in the UI using realistic hiring prompts.
- Inspect pass rate, grader reasoning, and activity-map evidence for at least one failing case.
- Use the findings to refine the agent before calling it release-ready. **Speaker note:** Connect the workflow back to the transcript-analysis architecture so quality improvement feels operational and repeatable.

Slide 84: Lab Time: Lab 25

- Start the optional developer branch while the main room stays on wrap-up, QA review, and rollout planning.
- Direct developers to the VS Code workflow only if they are caught up on the core path.
- Keep nondevelopers focused on evaluation findings, ROI framing, and next actions. **Speaker note:** This structure keeps the core workshop complete while still giving developers a relevant extension path.

Slide 85: Lab 25: VS Code workflow

- Optionally clone and sync the agent with the VS Code extension using the supported GA workflow.
- Validate that local edits flow back into Copilot Studio without changing the workshop baseline or bypassing in-product controls.

- Use this branch only for developers while the main room closes on evaluation, ROI, and rollout actions. **Speaker note:** End by making the developer add-on feel optional, useful, and safely separated from the core path.